

► Processes

- What is a process?
- Structure of a process.
- Process state.
- Schedulers and Scheduling Queues.
- Context switching.
- Process creation and termination.
- Inter Process Communication.

Dr. Manmath N. Sahoo
Dept. of CSE, NIT Rourkela

Process

- ▶ Process is a program that has initiated its execution.
- ▶ A program is a **passive** entity; whereas a process is an **active** entity.

Process structure

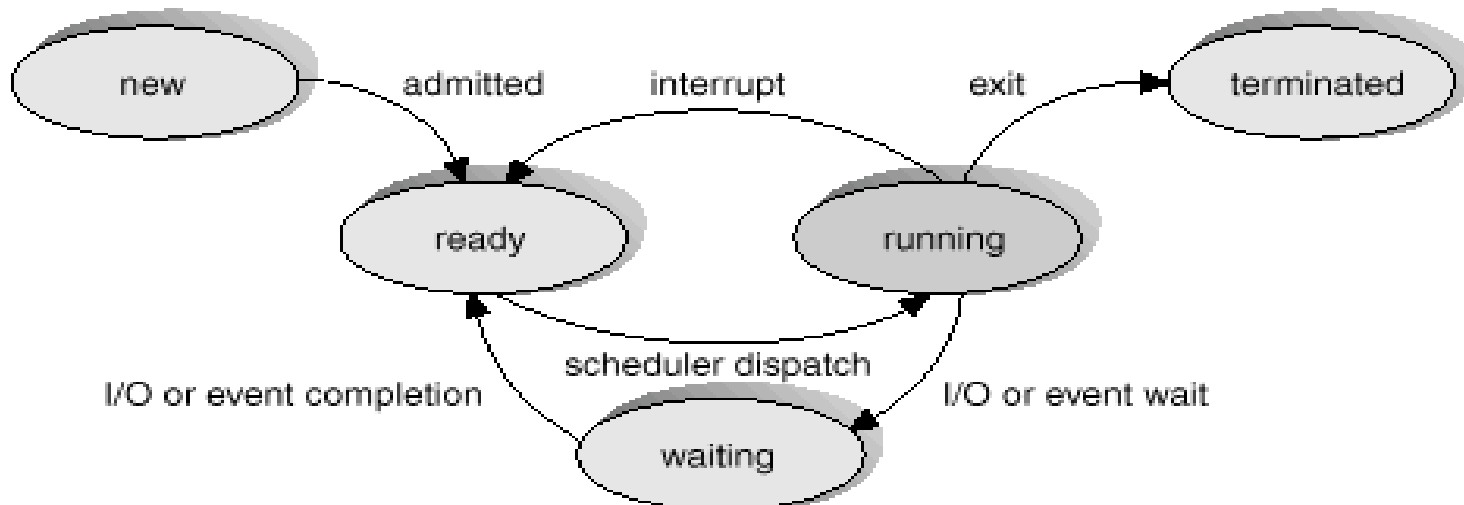
A process includes:

- ▶ program section - contains a copy of the machine code instructions
- ▶ user data section - to hold variable data values
- ▶ system data section - to hold
 - ▶ process context info when process interrupted - program counter, processor registers, etc and
 - ▶ system information about resources allocated to process, etc.
- ▶ user and system data sections obviously may be different from one run of a process to another

Process State

As a process executes, it changes state.

- ▶ **new**: The process is created by OS
- ▶ **ready**: The process is waiting to be assigned to a processor.
- ▶ **running**: Instructions are being executed on processor
- ▶ **waiting**: The process is waiting for some event to occur
- ▶ **terminated**: The process has finished execution.



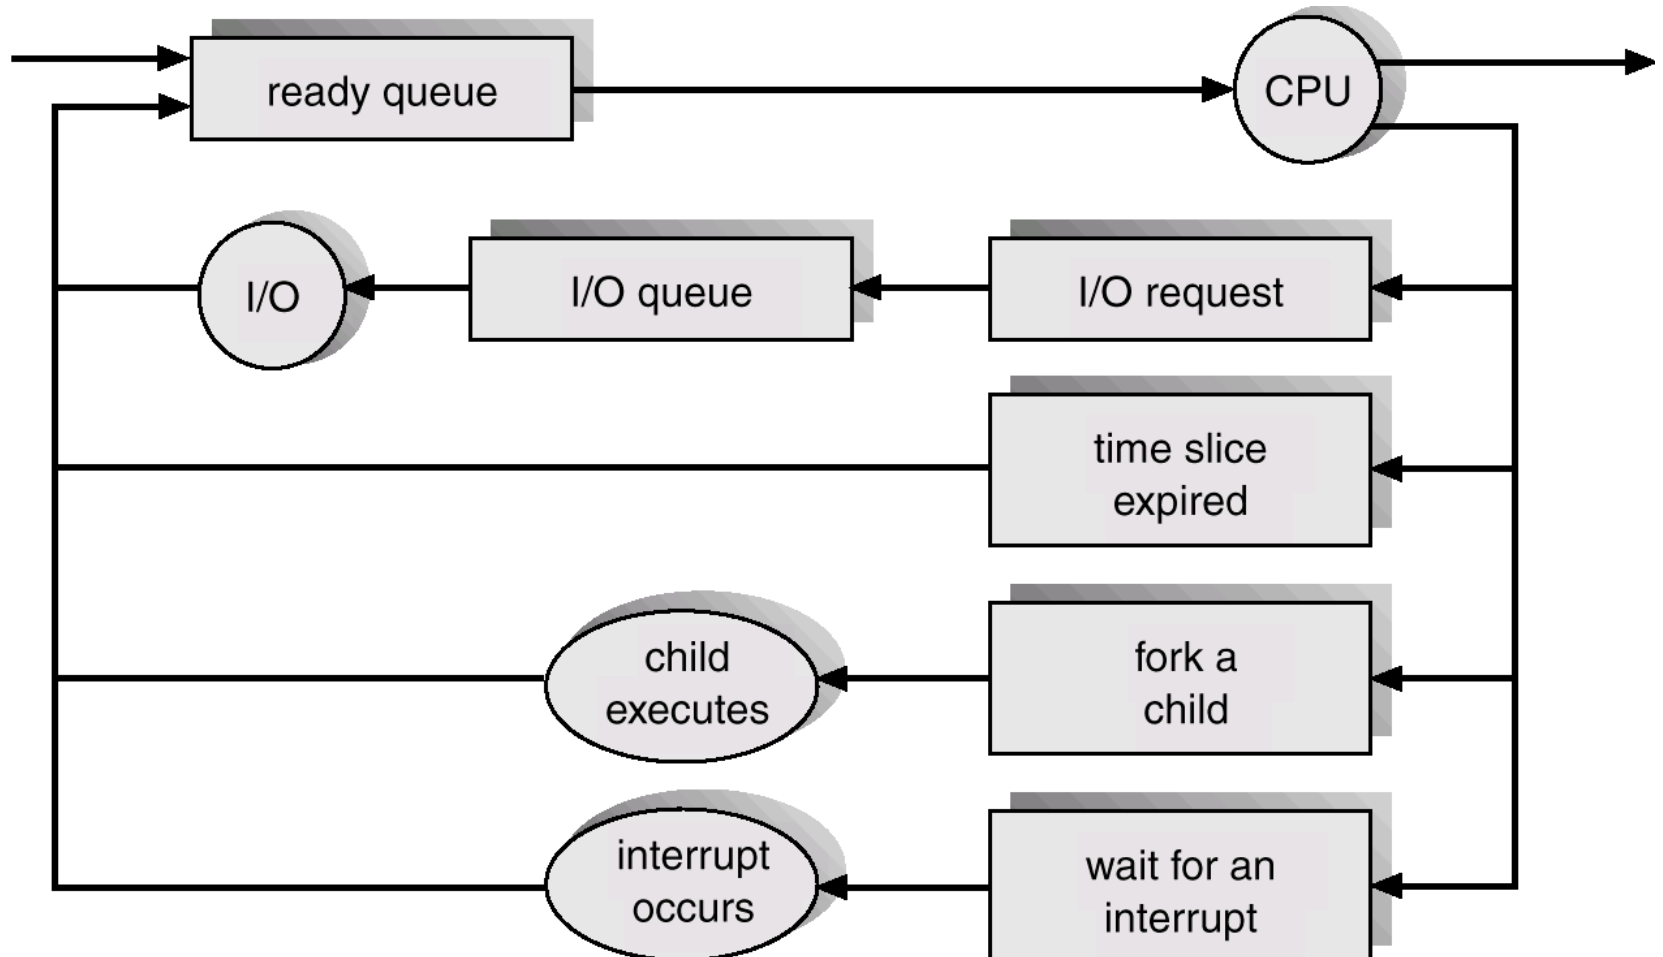
Process Control Block (PCB)

- ▶ PCB - data structure that contains Information associated with each process:
 - ▶ Process ID
 - ▶ Process State
 - ▶ Program Counter
 - ▶ Register Contents
 - ▶ Starting address of the process
 - ▶ Size of the process
 - ▶ Pointer to Child Process
 - ▶ Pointer to Next PCB (Sibling Pointer)
 - ▶ Resource Pointer

Process Scheduling Queues

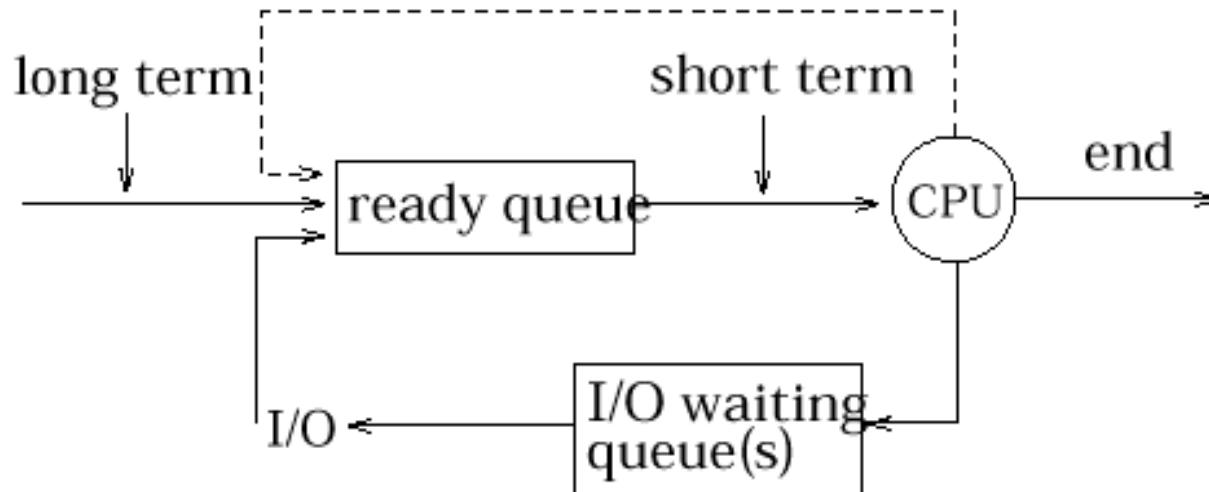
- ▶ **Job queue** – Set of all processes just created by LTS (new state processes).
- ▶ **Ready queue** – set of all processes that are in main memory, ready to execute (ready state processes).
- ▶ **Device queues** – set of processes waiting to use an I/O device (waiting state processes).

Queueing Diagram



Schedulers

- ▶ **Long-term scheduler (LTS)** (or job scheduler) – selects which processes should be brought into the ready queue.
- ▶ **Short-term scheduler (STS)** (or CPU scheduler) – selects which process should be executed next on CPU.



Schedulers

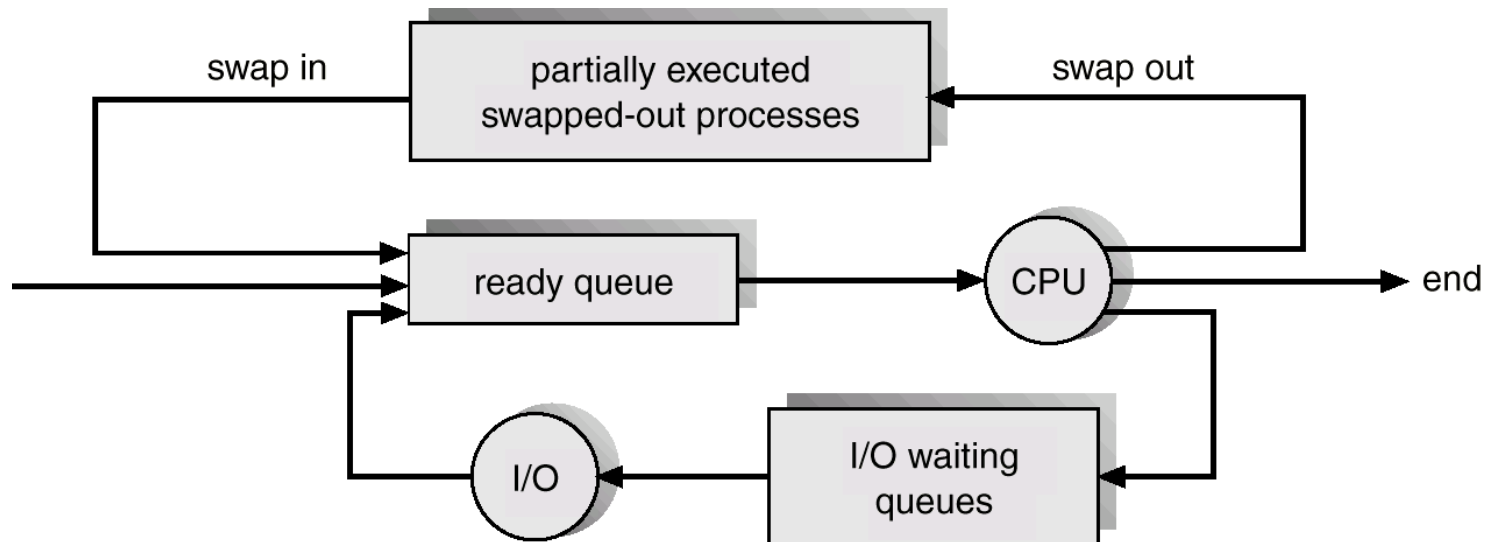
- ▶ Short-term scheduler is invoked very frequently (milliseconds) -> (must be fast).
- ▶ Long-term scheduler is invoked very infrequently (seconds, minutes) -> (may be slow).
- ▶ The long-term scheduler controls the degree of multiprogramming - this is the number of processes in the system that can be scheduled onto the CPU I.e. number of processes competing for CPU as a resource.

Schedulers

- ▶ Long-term schedule should choose a good process mix; otherwise it may lead to either low CPU utilization (if I/O bound processes are selected only) or low Device utilization (if CPU bound processes are selected only).
- ▶ I/O-bound process – have small amount of computation before it needs to do some I/O; many I/O requests -Typical interactive user programs.
- ▶ CPU-bound process – large amount of computation before it needs to do some I/O; few I/O requests - Programs that require sustained periods of calculation e.g. modelling applications

Schedulers

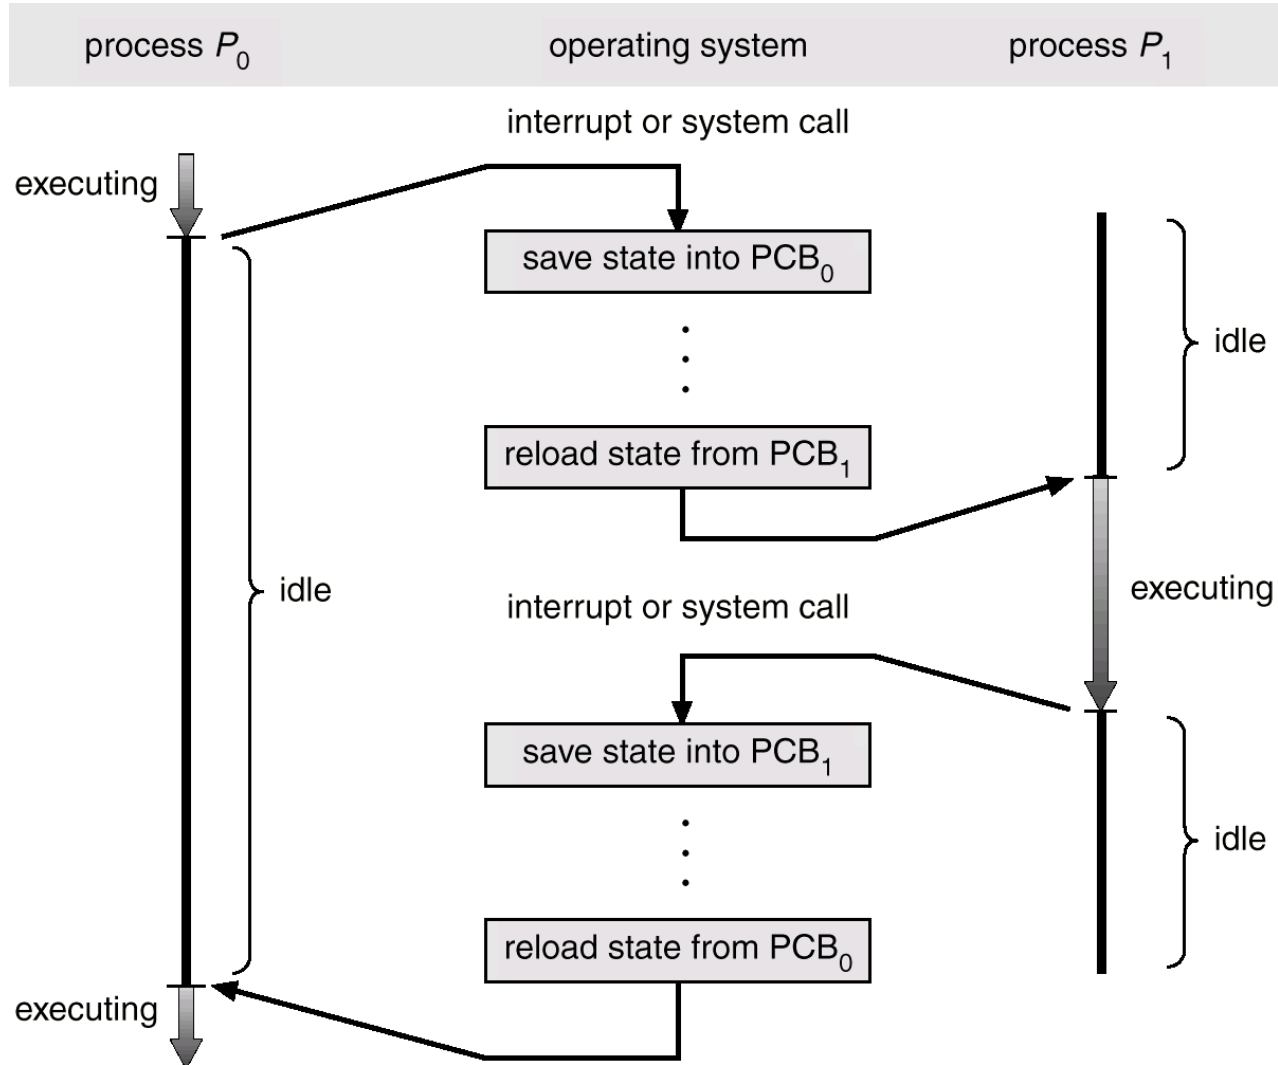
- ▶ Some OS have a light LTS; and a **Medium Term Scheduler (MTS)**.
- ▶ If degree of multiprogramming becomes high then total context switch overhead will be more.
- ▶ MTS can temporarily swap some processes out to disk and bring them back later.



Context Switch

- ▶ When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.
- ▶ Context-switch time is overhead; the system does no useful work while switching.

CPU switched between processes



Process Creation

- ▶ When a process is created it is created by another process known as the parent process - new process created is the child process of the parent. New process may in turn create other processes, forming a family tree of processes.
- ▶ There must be a first process (**init** process) - this is process that is created at bootup of the OS and is part of OS.

Process Creation (Cont.)

- ▶ Resource sharing between parent and child processes - options include:
 - ▶ Parent and children share all resources, or
 - ▶ Children share subset of parent's resources, or
 - ▶ Parent and child share no resources.
- ▶ Execution options:
 - ▶ Parent and children execute concurrently, or
 - ▶ Parent waits until children terminate.

Process Termination

- ▶ Process executes last statement and asks the operating system to delete it (**exit** system call). Then
 - ▶ Return data from child to parent.
 - ▶ Process' resources are deallocated by operating system.
- ▶ Parent may terminate execution of children processes (**abort**). Reasons for this could be:
 - ▶ Task assigned to child is no longer required.
 - ▶ Parent is exiting and O/S does not allow children to continue if their parent terminates – **cascaded termination**.

Cooperating Processes

- ▶ Process is **independent** if it cannot legally affect or be affected by the execution of another process.
- ▶ Process is **cooperating** if it can legally affect or be affected by the execution of another process.

Why Cooperating Processes?

- ▶ Processes may need to share data
 - ▶ More than one process reading/writing the same data (a shared file, a database record,...)
 - ▶ Output of one process being used by another
- ▶ Ordering executions of multiple processes may be needed to ensure correctness
 - ▶ Process X should not do something before process Y does something etc.
 - ▶ Need mechanisms to pass control signals between processes
- ▶ Modularity - dividing system functions into separate processes that then talk to each other
- ▶ Computation speed-up - but only if >1 processor

Inter-Process Communication (IPC)

Fundamental types of IPC

- ▶ Message passing


- ▶ P and Q exchange messages

- ▶ Shared memory

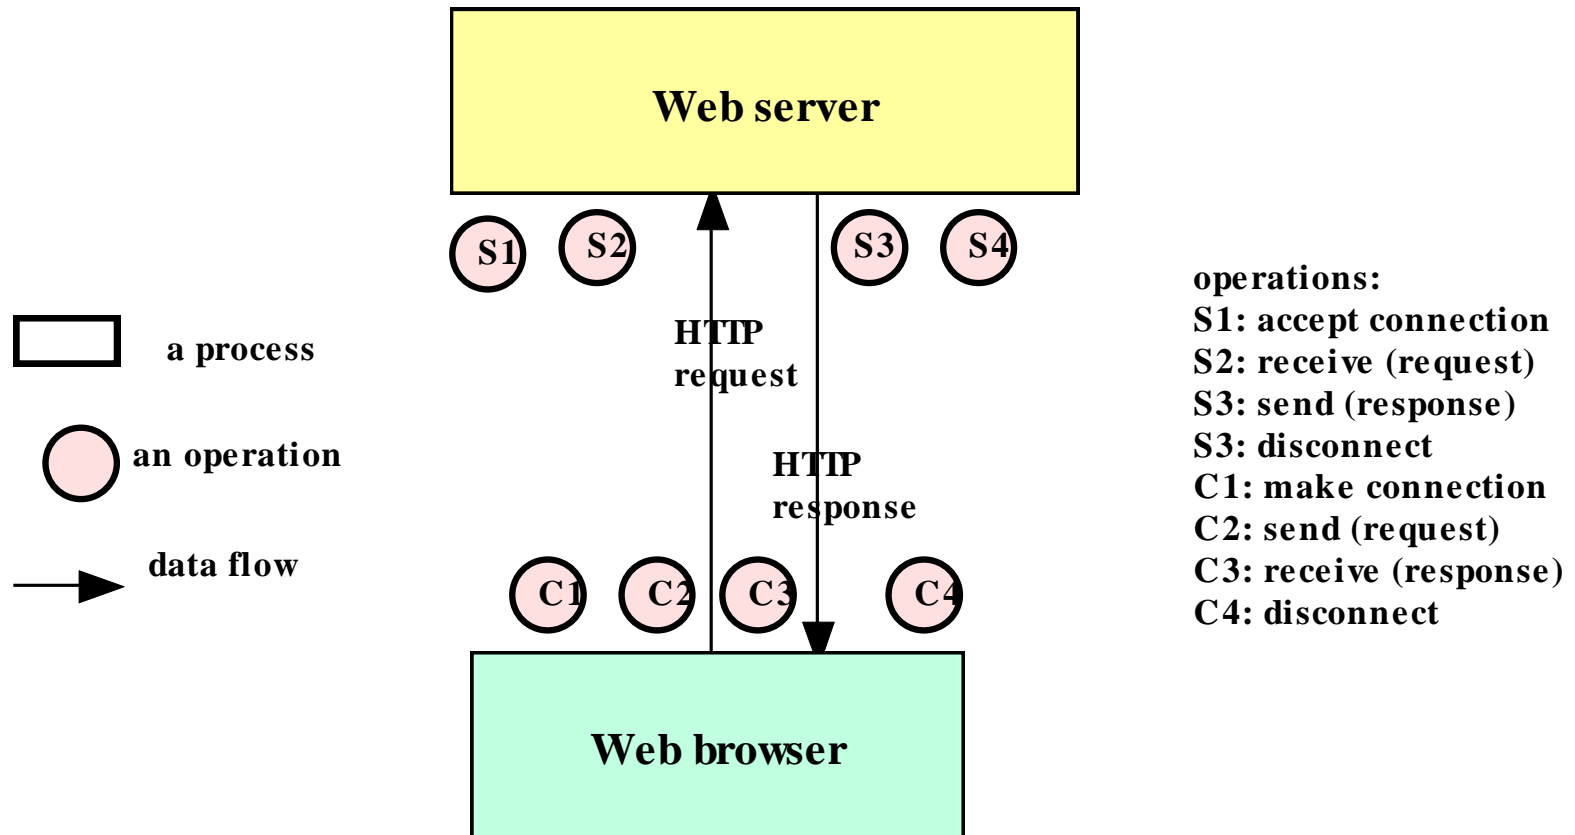
- ▶ P writes into a shared location, Q reads from it and vice-versa

Inter-Process Communication (IPC)

(a) Message Passing (b) Shared Memory



Message Passing System



Message Passing System

- ▶ Direct vs. Indirect
- ▶ Symmetric vs. Asymmetric
- ▶ Synchronous vs. Asynchronous

Message Passing System: Direct

Primitives:

- ▶ **Connect (sender address, receiver address)**, for connection-oriented communication.
- ▶ **Send ([receiver], message)** e.g., Send(P, msg)
- ▶ **Receive ([sender], message storage object)** e.g., Receive(Q,msg)
- ▶ **Disconnect (connection identifier)**, for connection-oriented communication.

Message Passing System: Direct

- ▶ **Symmetric** – Both the sender and receiver have to explicitly name each other.
- ▶ **Asymmetric** – Only the sender has to name the receiver, but the receiver doesn't need to name the sender.
 - ▶ Facilitates one-to-many communication.
 - ▶ `Send(P,msg);` - send message to P.
 - ▶ `Receive(id,msg);` - receive message from any process. “**id**” will be set to the name of the process with which communication has taken place.

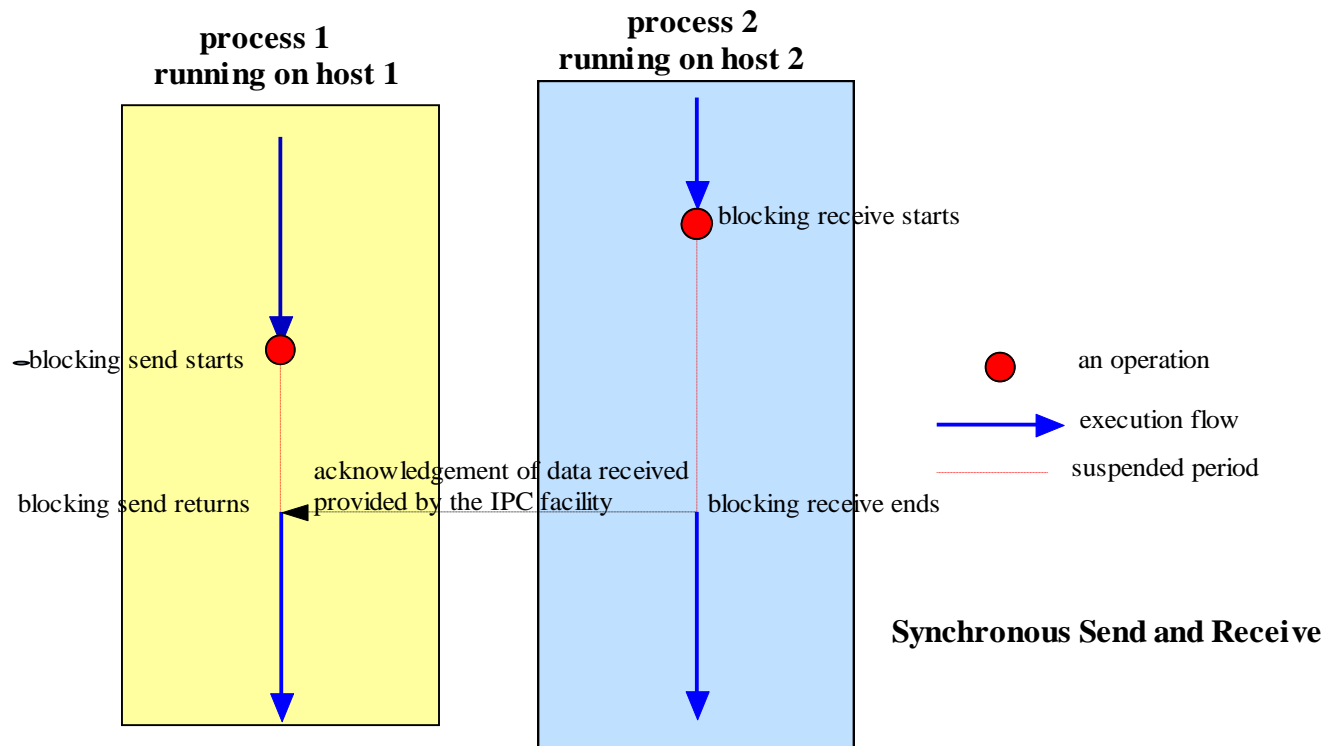
Message Passing System: Indirect

- ▶ The messages are sent and received to/from **mailboxes**.
- ▶ Each mailbox has a unique identifier.
- ▶ Two processes can communicate with each other iff they have a shared mailbox.

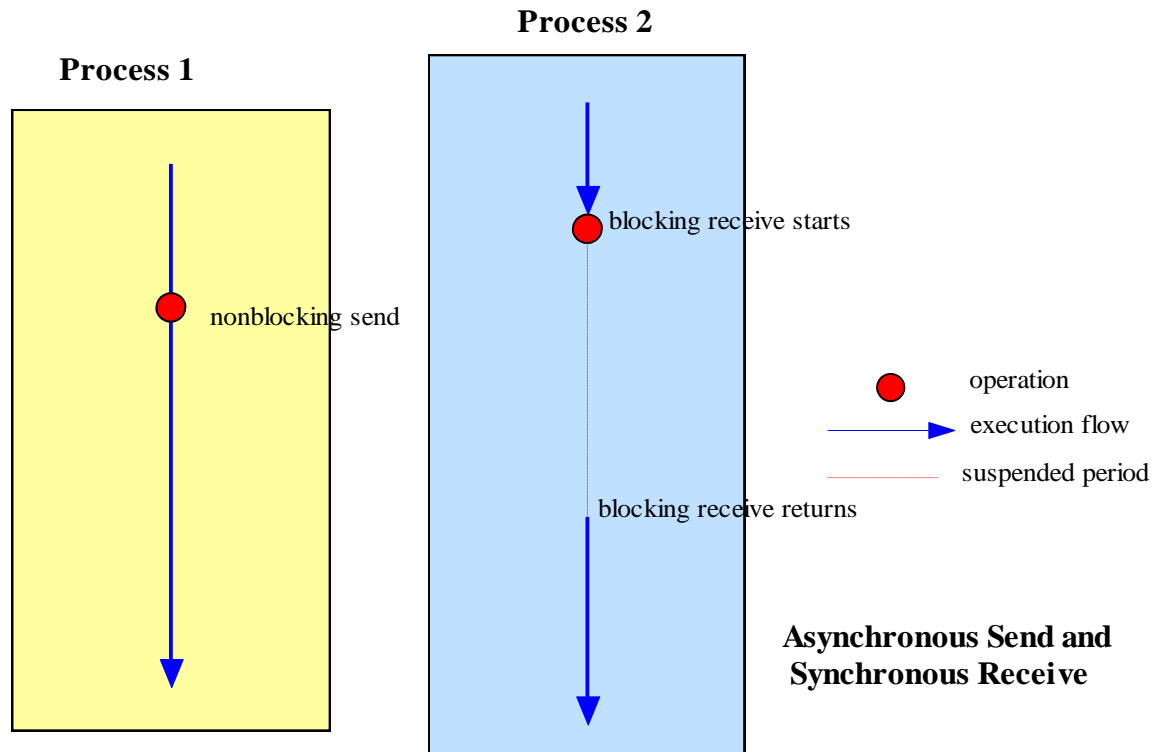
Message Passing System: Synchronous Vs. Asynchronous

- ▶ Synchronous – After initiating a send/receive operation, the process waits for the acknowledgement (the process is blocked)
- ▶ Asynchronous – After initiating a send/receive operation, the process continues its execution

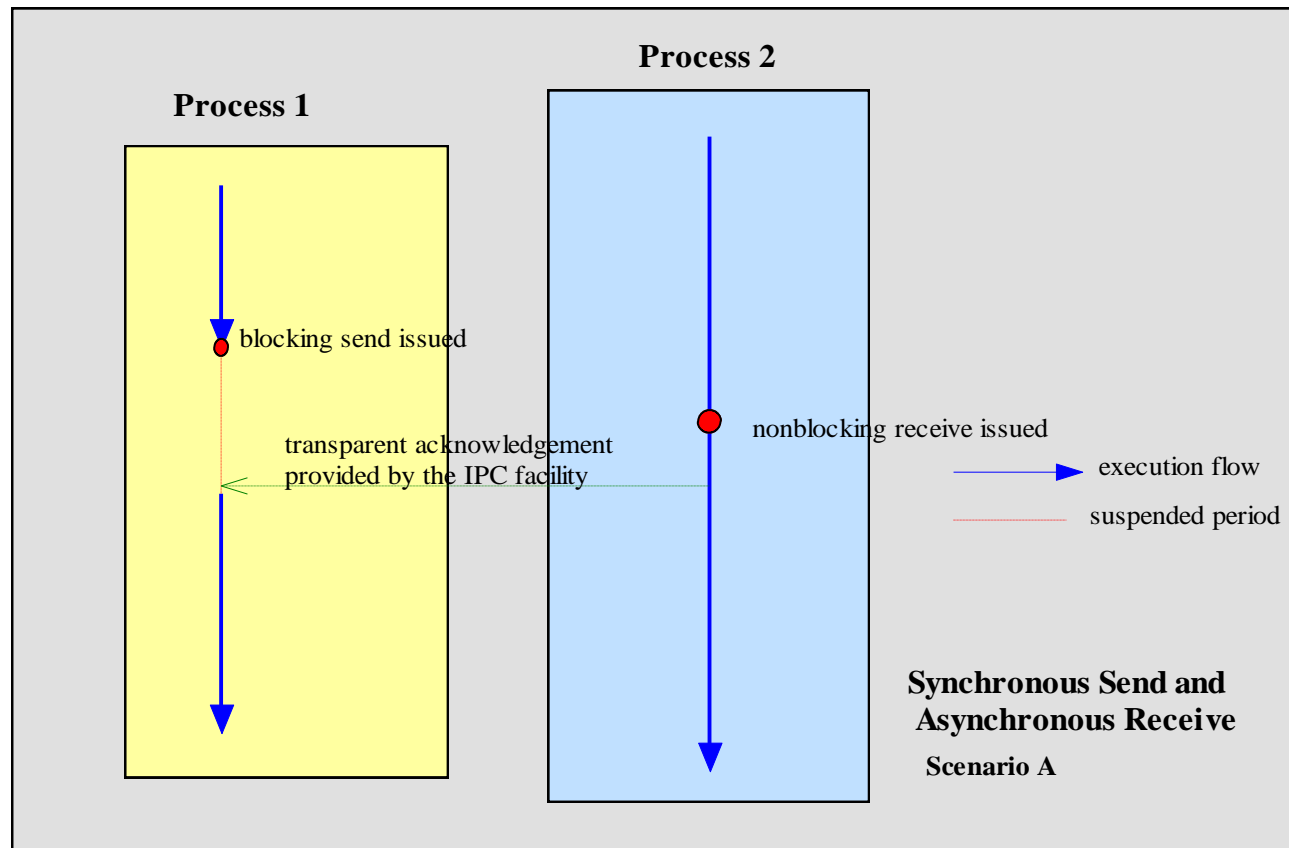
Synchronous send and receive



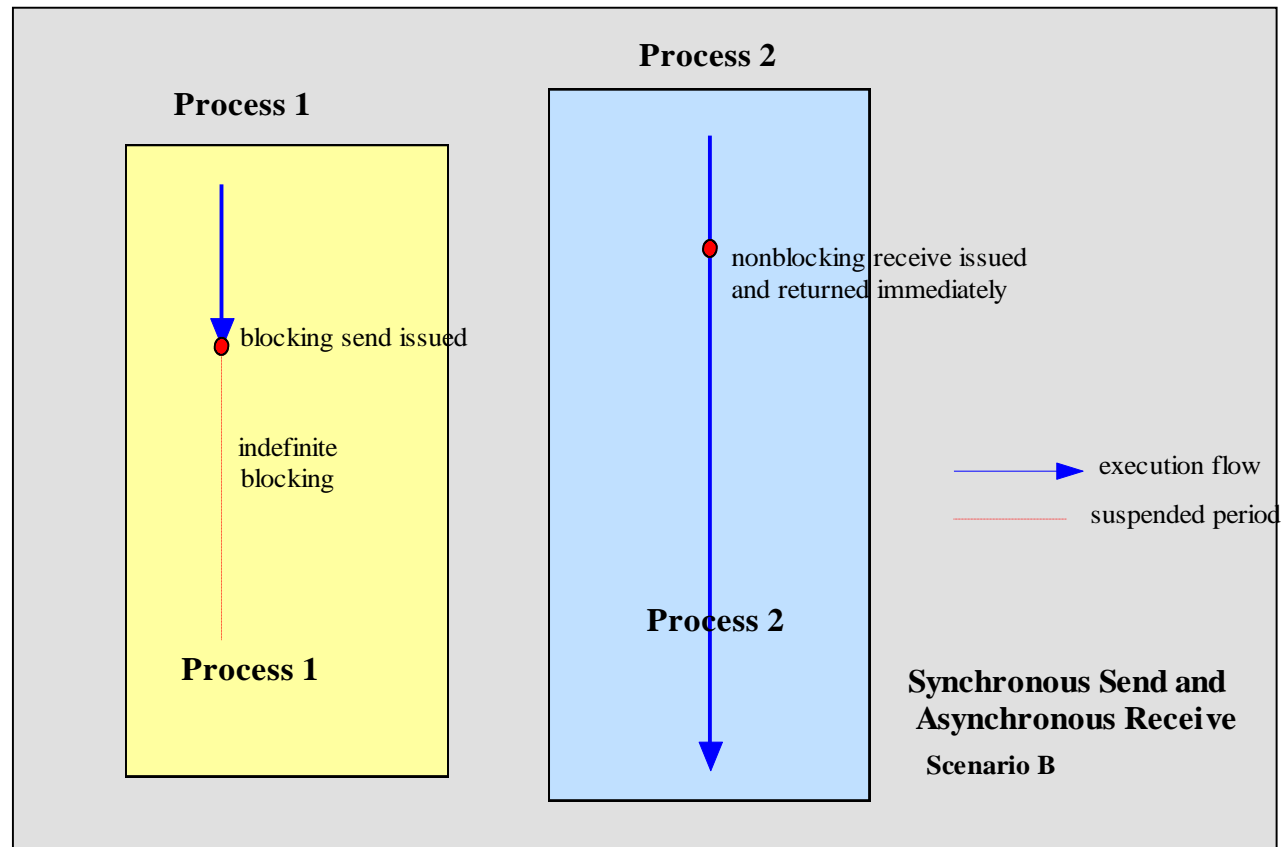
Asynchronous send and synchronous receive



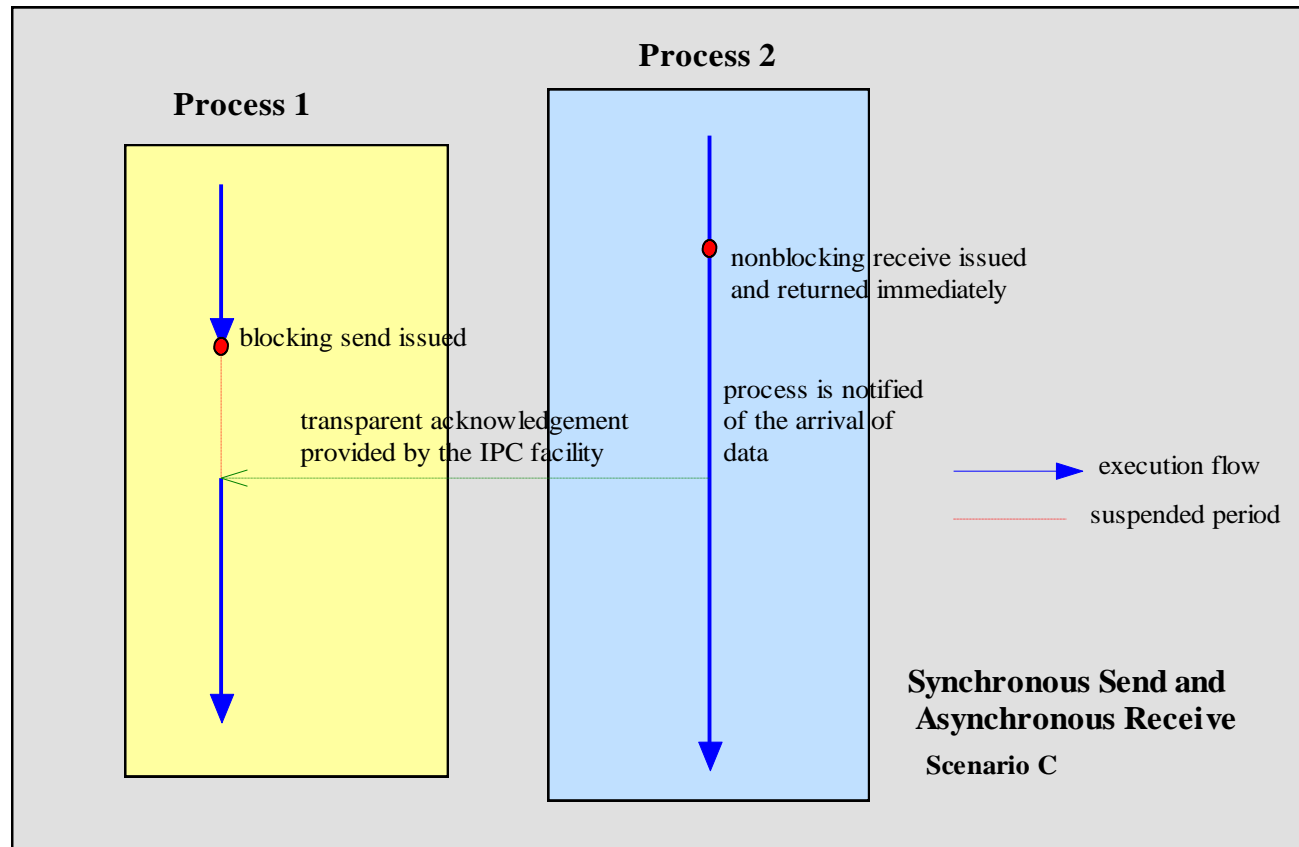
Synchronous send and Async. Receive - A



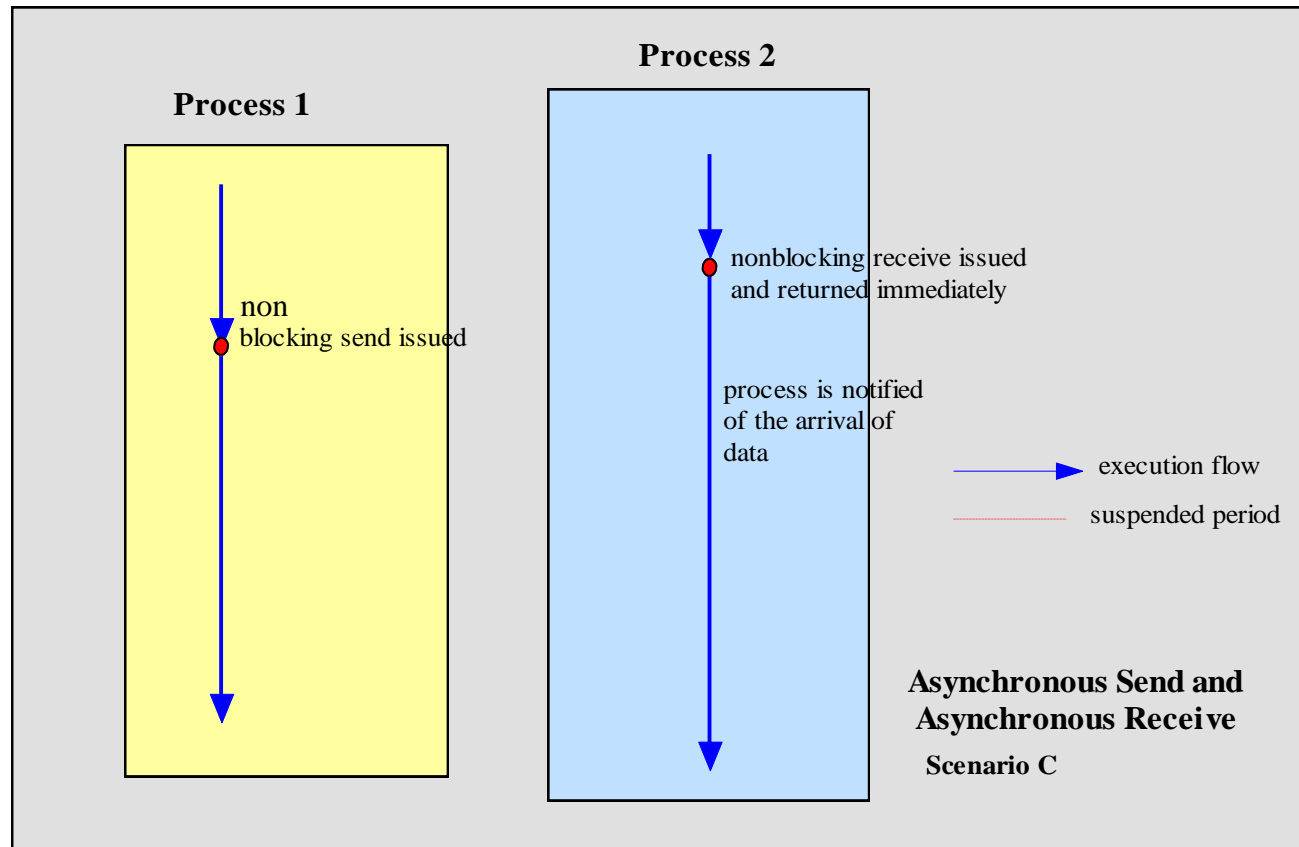
Synchronous send and Async. Receive - B



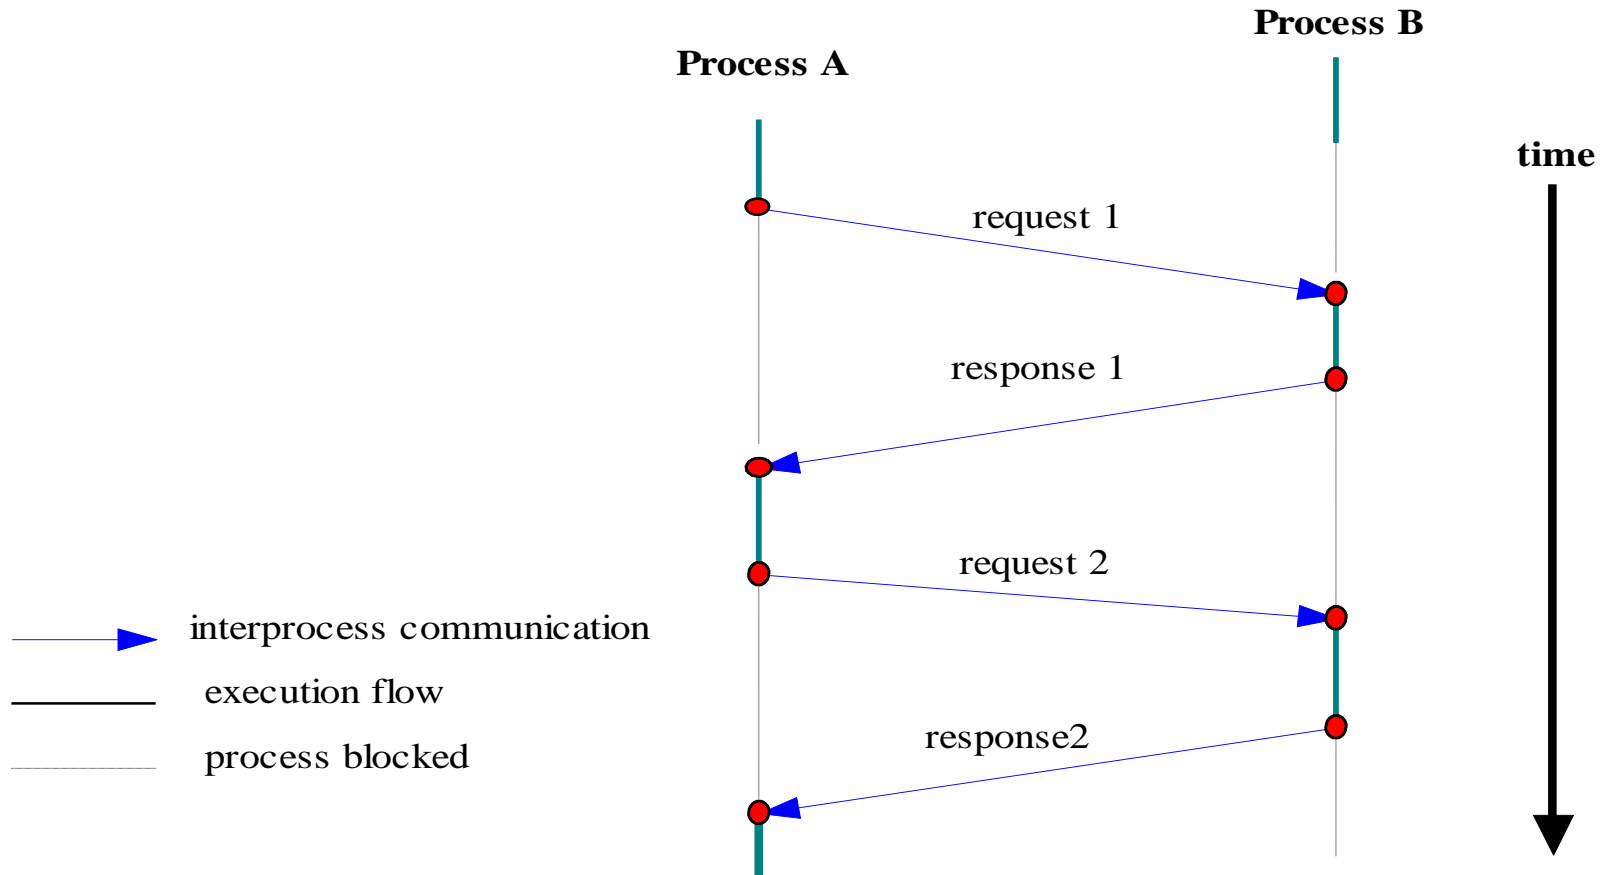
Synchronous send and Async. Receive - C



Asynchronous send and Asynchronous receive - C



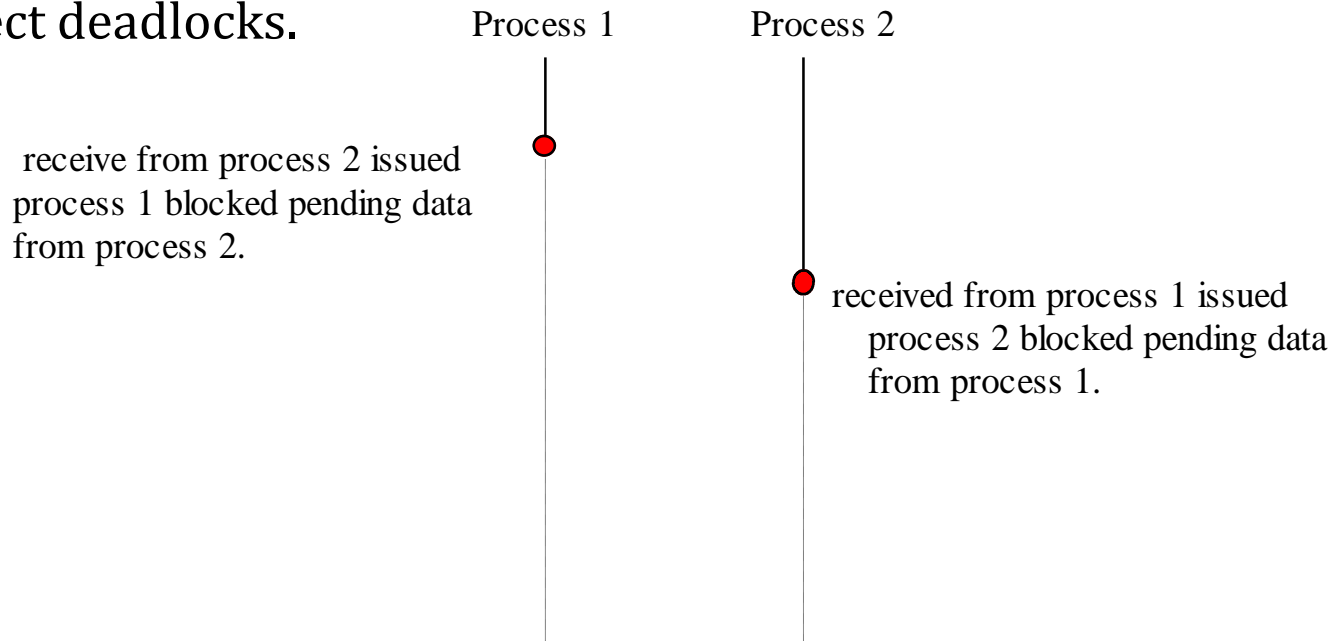
Event diagram



Event diagram for a protocol

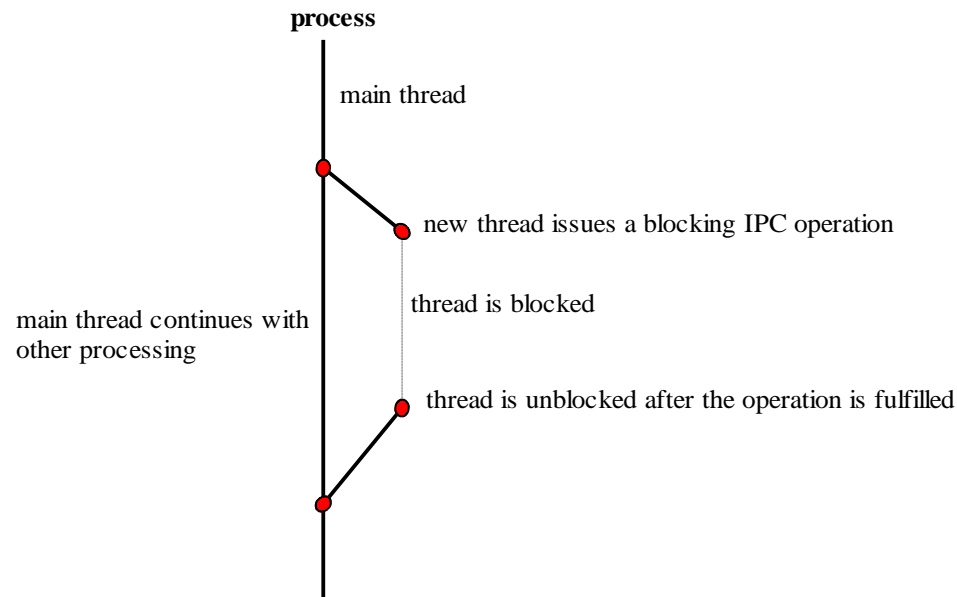
Blocking, deadlock, and timeouts

- ▶ Blocking operations issued in the wrong sequence can cause deadlocks.
- ▶ Deadlocks should be avoided. Alternatively, timeout can be used to detect deadlocks.



Using threads for asynchronous IPC

- If only blocking operation is provided for send and/or receive, then it is the programmer's responsibility to use child processes or threads if asynchronous operations are desired.



Buffering during message passing

▶ Zero capacity

- ▶ No waiting queue/buffer
- ▶ Sender must wait for receiver, after generating one item

▶ Bounded capacity

- ▶ Buffer has a fixed length N .

▶ Unbounded capacity

- ▶ No practical limit on the size of the buffer

Process: Exception Conditions

▶ Process Termination

- ▶ Sender Terminates: If the receiver has initiated a blocking receive then the OS either terminates the receiver or notifies it that Sender has terminated.
- ▶ Receiver Terminates: If the sender has initiated a blocking send then the OS either terminates the Sender or notifies it that Receiver has terminated.

▶ Lost Messages

- ▶ The OS is responsible for detecting this event and for resending the message.
- ▶ The Sender is responsible for detecting this event and for resending it.
- ▶ The OS detects the event and notifies the Sender about it.

▶ Scrambled Messages

- ▶ Due to noise in the channel
- ▶ Handled in the similar ways as lost messages.