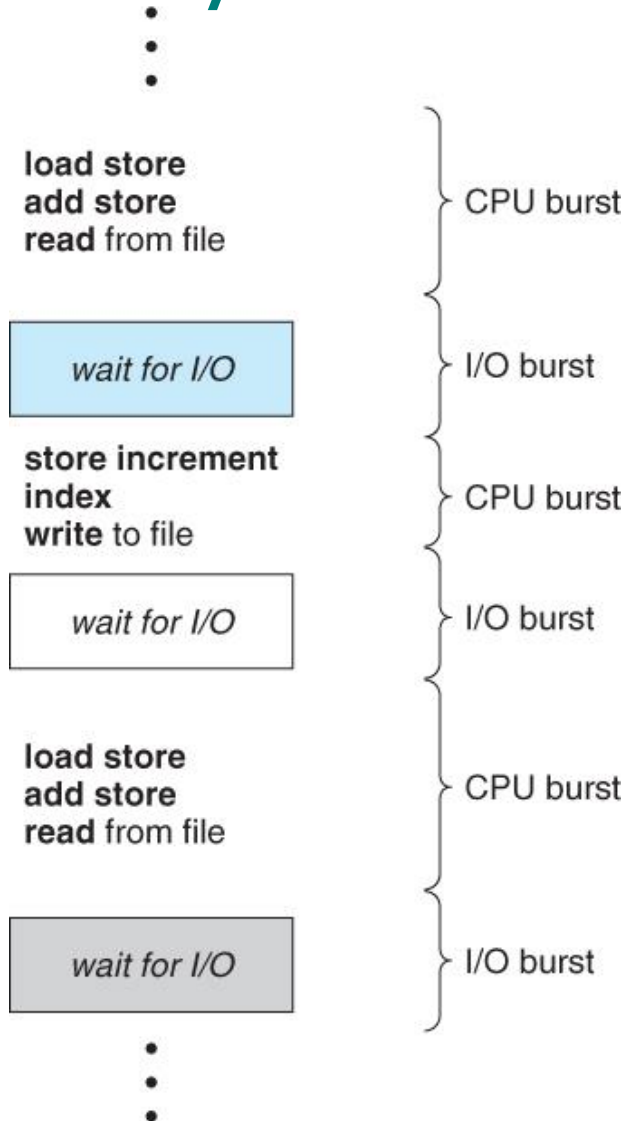


► CPU Scheduling

Dr. Manmath N. Sahoo
Dept. of CSE, NIT Rourkela

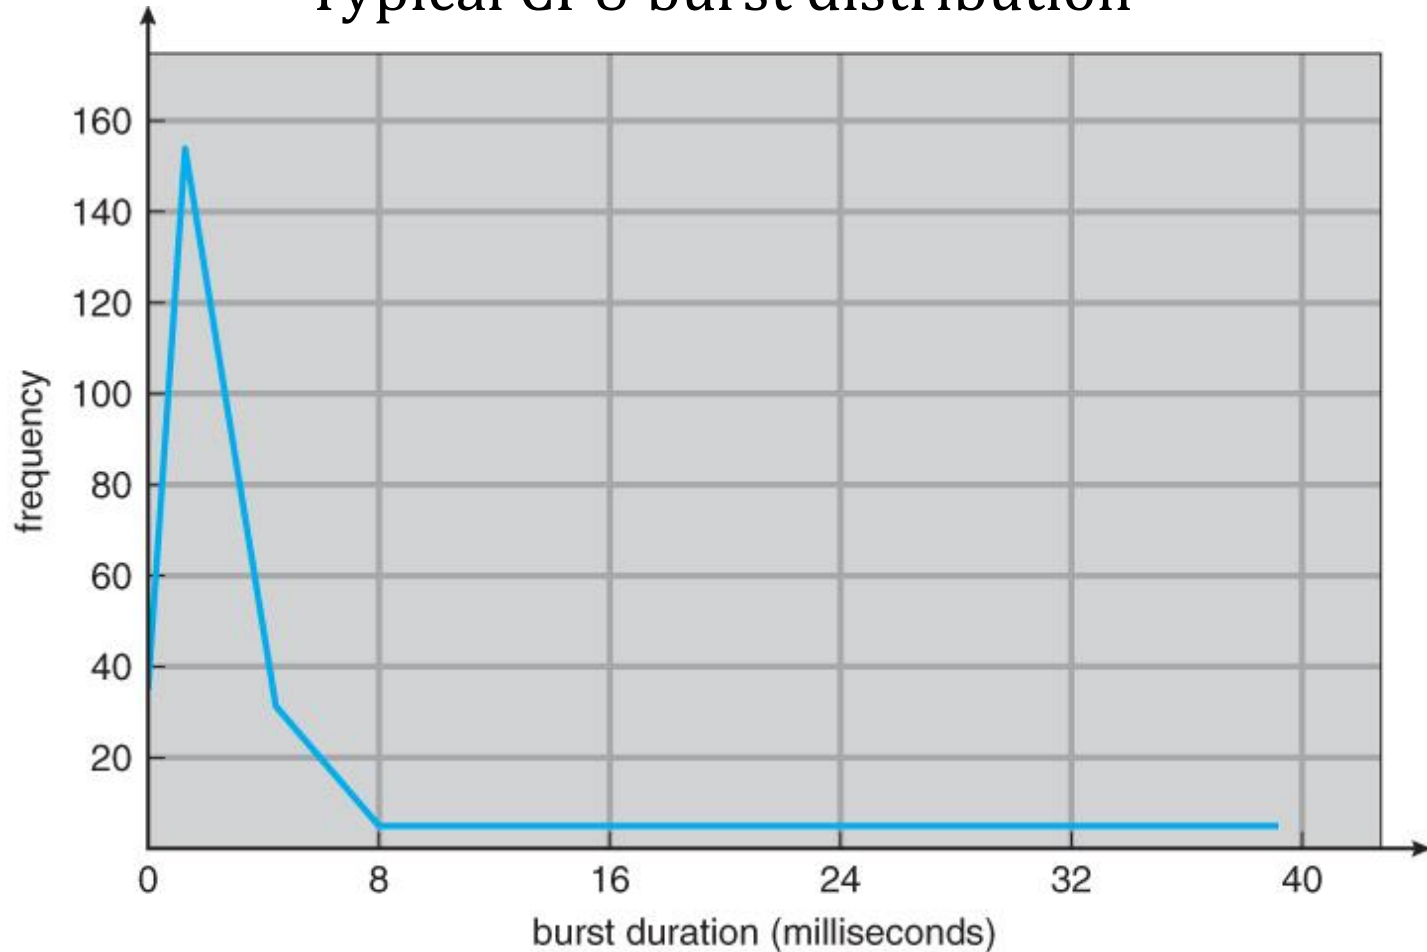
CPU-I/O Burst Cycle



- ▶ CPU-I/O Burst Cycle – Process execution consists of a cycle of CPU execution and I/O wait.
- ▶ CPU burst is length of time process needs to use CPU before it next makes a I/O request.
- ▶ I/O burst is the length of time process spends waiting for I/O to complete.

Histogram of CPU-burst Times

Typical CPU burst distribution



CPU Scheduler

- ▶ Allocates CPU to one of processes that are ready to execute (in ready queue)
- ▶ CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state (e.g. I/O request)
 2. Terminates
 3. Switches from waiting to ready (e.g. I/O completion)
 4. Switches from running to ready state (e.g. priority based interruption)
- ▶ Scheduling only when 1 and 2 happen is **nonpreemptive** - process keeps CPU until it voluntarily releases it
- ▶ Scheduling also when 3 & 4 happen is **preemptive** - CPU can be taken away from process by OS

Dispatcher

- ▶ **Dispatcher** gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ▶ switching context
 - ▶ switching to user mode (if CPU was executing the previous process in kernel mode and preempted)
 - ▶ jumping to the proper location in the user program to restart that program (i.e., last action is to set program counter)
- ▶ **Dispatch latency** – time it takes for the dispatcher to switch between processes and start new one running

Scheduling Criteria

- ▶ **CPU utilization:** i.e. CPU usage – want to maximize
- ▶ **Throughput:** number of processes that complete their execution per time unit – want to maximize
- ▶ **Turnaround time:** amount of time to execute a particular process – want to minimize
- ▶ **Waiting time:** amount of time a process has been waiting in the ready queue – want to minimize
- ▶ **Response time:** amount of time it takes from when a job was submitted until it initiates its first response (output) – want to minimize

First-Come, First-Served Scheduling

- Schedule: order of arrival of process in ready queue

Example:

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3


- Suppose that the processes arrive in the order: P_1, P_2, P_3 . The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling

- ▶ Suppose that the processes arrive in the order P_2, P_3, P_1 .
- ▶ The Gantt chart for the schedule is:



- ▶ Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- ▶ Average waiting time: $(6 + 0 + 3)/3 = 3$

FCFS Scheduling

- + Simple to implement by FIFO queue.
- + Code is easy to write and understand
- Average waiting time is not minimal
- Not suitable for time-sharing systems
- Convoy effect: short I/O intensive processes behind long CPU intensive process

► **NOTE: FCFS Scheduling is non-preemptive**


Shortest-Job-First (SJF) Scheduling

- ▶ Schedule: order of the CPU burst of the processes.
- ▶ Two schemes:
 - ▶ Non-Preemptive SJF / Shortest Job Next (SJN)
 - ▶ once CPU given to the process it cannot be preempted until completes its CPU burst.
 - ▶ Preemptive SJF / Shortest Remaining Time First (SRTF)
 - ▶ if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.

Non-Preemptive SJF (SJN)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	7
P_2	2	4
P_3	4	1
P_4	5	4

► Gantt Chart:




► Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Preemptive SJF (SRTF)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	7
P_2	2	4
P_3	4	1
P_4	5	4

► Gantt Chart:



► Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
<i>P1</i>	0	16
<i>P2</i>	2	8
<i>P3</i>	4	12
<i>P4</i>	6	2
<i>P5</i>	8	4

- ▶ Draw Gantt chart and find the average waiting time in each of the following cases:
 - ▶ SRTF Scheduling
 - ▶ SJN Scheduling
 - ▶ FCFS Scheduling

Answer to Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
<i>P1</i>	0	16
<i>P2</i>	2	8
<i>P3</i>	4	12
<i>P4</i>	6	2
<i>P5</i>	8	4

► SRTF



► Average waiting time = $(26+6+12+0+0)/5$
 $= 8.8$

Answer to Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
<i>P1</i>	0	16
<i>P2</i>	2	8
<i>P3</i>	4	12
<i>P4</i>	6	2
<i>P5</i>	8	4

► SJN



► Average waiting time = $(0+20+26+10+10)/5$
 $= 13.2$

Answer to Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
<i>P1</i>	0	16
<i>P2</i>	2	8
<i>P3</i>	4	12
<i>P4</i>	6	2
<i>P5</i>	8	4

► FCFS



► Average waiting time = $(0+14+20+30+30)/5$
 $= 18.8$

Answer to Example

Scheduling Algorithm	Average Waiting Time
FCFS	18.8
SJN	13.2
SRTF	8.8

Shortest Job First

- + SRTF is Optimal: Minimum Average waiting time.
- Longer processes may be starved.
- Difficult to know the CPU burst of processes before their execution

Determining Length of Next CPU Burst

- ▶ Can only estimate the length.
- ▶ Can be done by using the length of previous CPU bursts, using exponential averaging (decaying average).

Determining Length of Next CPU Burst

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n.$

Determining Length of Next CPU Burst

- ▶ $\alpha = 0, \tau_{n+1} = \tau_n$
 - ▶ last CPU burst does not count - only longer term history
- ▶ $\alpha = 1, \tau_{n+1} = t_n$
 - ▶ Only the actual last CPU burst counts.
- ▶ If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

- ▶ Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor.


Round Robin (RR) Scheduling

- ▶ Each process gets a **time quantum**, usually 10-100 milliseconds.
- ▶ After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ▶ If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time.
- ▶ No process waits more than $(n-1)q$ time units.

Example: RR with $q=20$

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

► The Gantt chart is:



► Typically, higher average turnaround than SJF, but better *response*.

Round Robin (RR) Scheduling

▶ FCFS + Preemption = RR

▶ Performance

▶ q large \Rightarrow behaves similar to FCFS

▶ q small \Rightarrow more context switch overhead

Modified RR Scheduling

- ▶ Based on variable time quantum
- ▶ A preempted process, on its next turn will use an increased time quantum
- ▶ RR Scheduling
 - ▶ P_1 : 300ms, $q=20$ ms, time per context switch=2ms
 - ▶ #of preemptions due to this process = 14
 - ▶ Preemption overhead = 28ms

Modified RR Scheduling

- ▶ Modified RR Scheduling – q is doubled in next turn.
- ▶ P_1 : 300ms, $q=20$ ms (initial)
- ▶ #of preemptions due to this process = 3.
($20+40+80+160=300$)
- ▶ Preemption overhead = 6ms

Priority Scheduling

- ▶ A priority number (integer) is associated with each process and the CPU is allocated to the process with the highest priority
- ▶ Some schemes use smallest integer \equiv highest priority.
- ▶ SJF is effectively priority scheduling where priority is defined on next CPU burst time.
- ▶ FCFS is effectively priority scheduling where priority is defined on arrival time.

Priority Scheduling

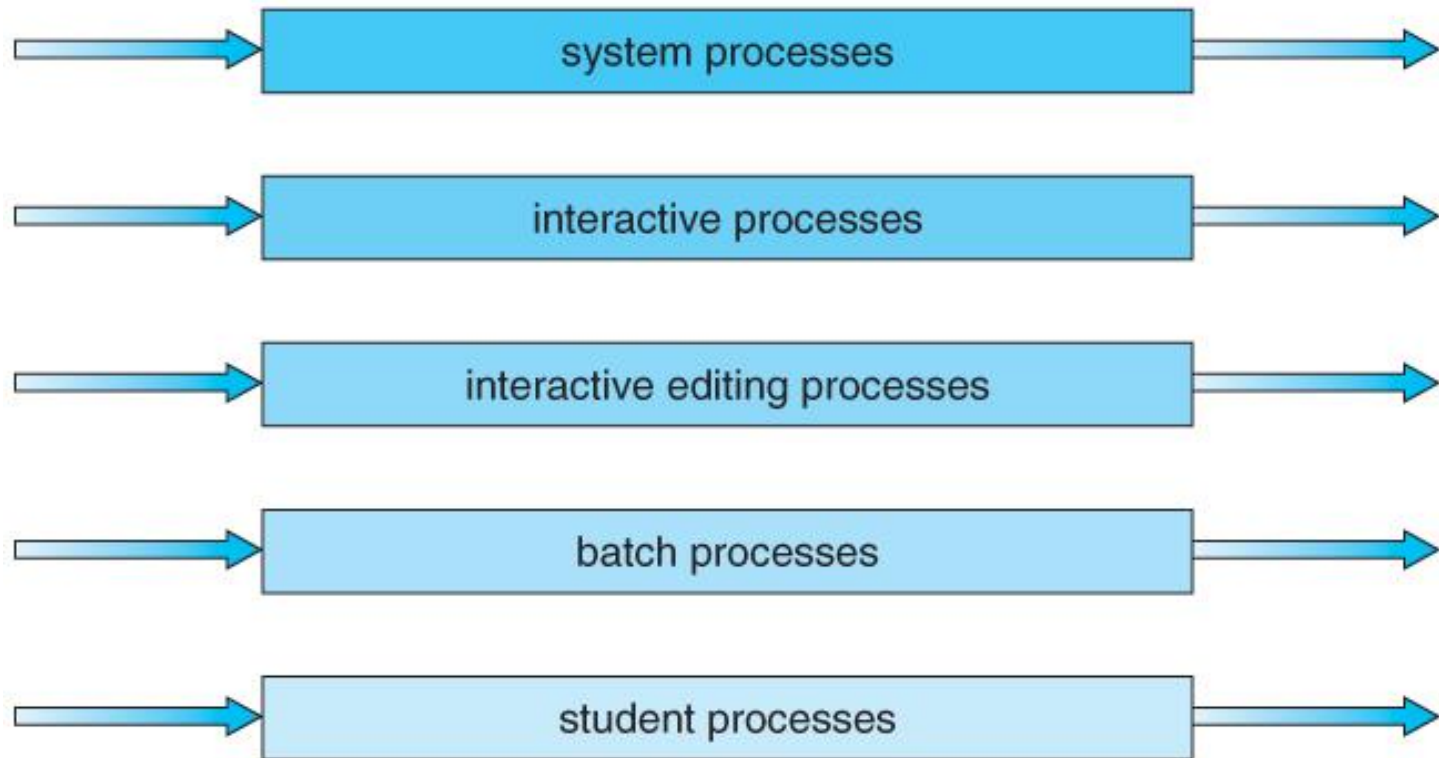
- ▶ Problem: **Starvation** – low priority processes may never execute.
- ▶ Solution: **Aging** – as time progresses increase the priority of the process.
 - ▶ e.g. in every 10ms increase the priority of process by 1

Multi level Queue Scheduling

- ▶ Ready queue is partitioned into separate queues.
- ▶ Each queue holds a particular category of processes.
- ▶ Each queue can implement whatever scheduling algorithm is most appropriate for that type of processes.

Multi level Queue Scheduling

highest priority



lowest priority

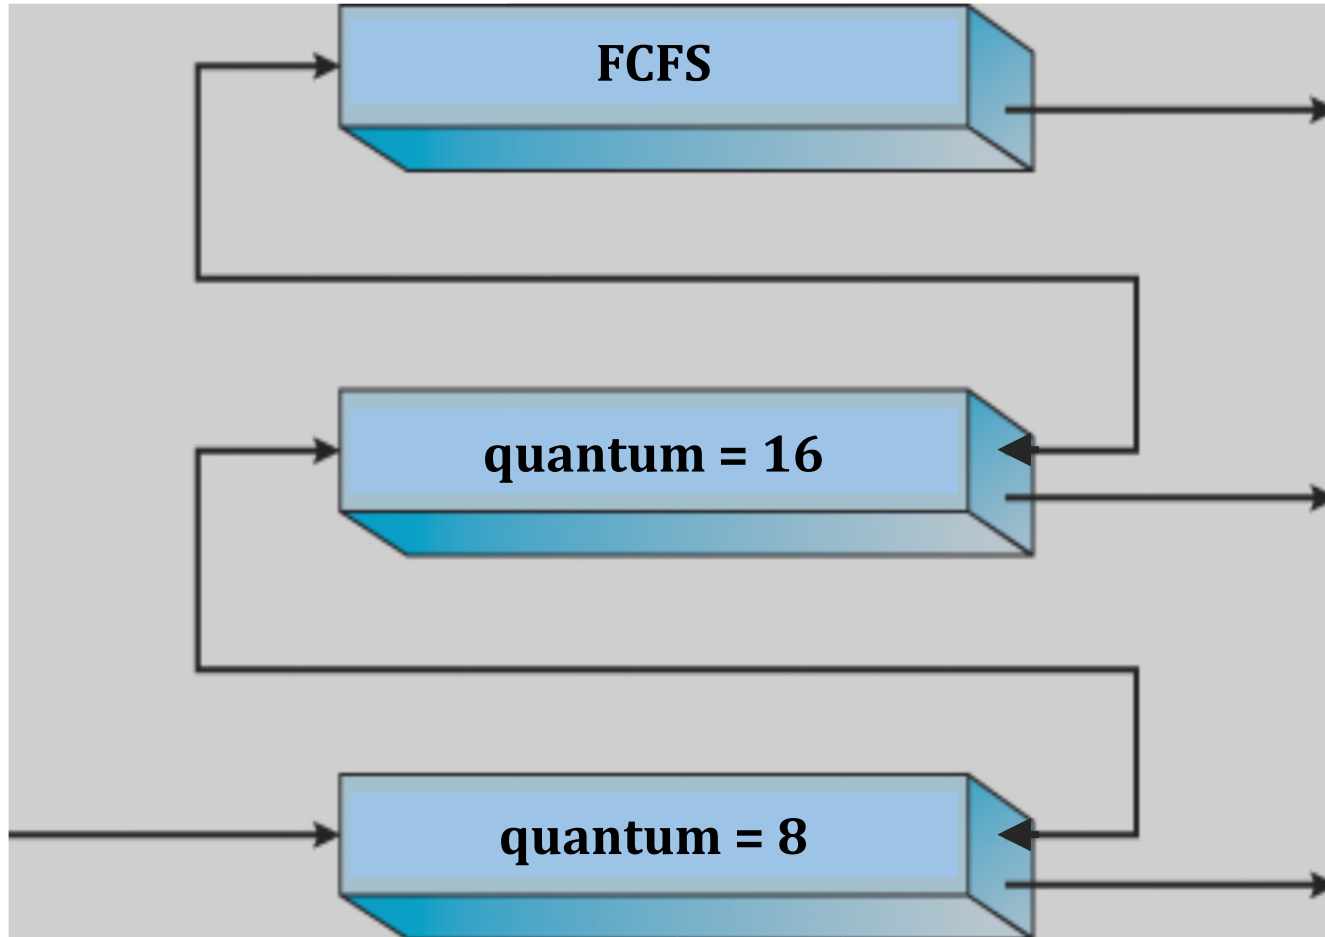
Multilevel Queue Scheduling

- ▶ Scheduling must be done between the queues.
- ▶ Priority scheduling; But possibility of starvation.
- ▶ Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes;

Multilevel Feedback Queue Scheduling

- ▶ A process can move between the various queues; **aging** can be implemented this way.
- ▶ Multilevel-feedback-queue scheduler defined by the following parameters:
 - ▶ number of queues
 - ▶ scheduling algorithms for each queue
 - ▶ method used to determine when to upgrade a process
 - ▶ method used to determine when to demote a process
 - ▶ method used to determine which queue a process will enter initially

Multilevel Feedback Queue Scheduling



Multi Processor Scheduling

- ▶ Separate queue for each processor
- ▶ Common ready queue

Multi Processor Scheduling: Separate queues

► 4 Processors: C1, C2, C3, C4.

► P1 goes to C1

► P2 goes to C2

► P3 goes to C3

► P4 goes to C4

► P5 goes to C1

► P6 goes to C2 ...

► Load balancing w.r.t size of the processes???

► C1 may be overloaded, while C3 is idle

Multi Processor Scheduling:

Common ready queue

► Symmetric Approach:

- Each processor examines the common ready queue and selects a process to execute.
- Several instances of STS will be running on different processors, whenever necessary
- Mutual exclusive access to the common queue???

► Asymmetric Approach

- One processor acts as scheduler (Master) for other processor (Slaves)