

# ► Deadlock

Dr. Manmath N. Sahoo  
Dept. of CSE, NIT Rourkela

# Shareable vs Non-shareable resources

## ▶ Sharable resources

- ▶ can be used by multiple processes during their execution.
- ▶ e.g., CPU, I/O bus.

## ▶ Non-sharable resources

- ▶ cannot be used by multiple processes during their execution.
- ▶ e.g., printer

# Static vs. Dynamic resource allocation

- ▶ **Static:** Allocate all the required resources prior to the execution of the process.
- ▶ **Dynamic:** Allocate only the initially required resources to start up the process; additional resources will be allocated dynamically.

Static	Dynamic
<ul style="list-style-type: none"><li>▶ Simpler to implement</li><li>▶ Process is guaranteed to complete in definite time</li><li>▶ Resource utilization less</li><li>▶ Deadlock will never occur</li></ul>	<ul style="list-style-type: none"><li>▶ Comparatively complex</li><li>▶ No certainty about the process completion</li><li>▶ Better resource utilization</li><li>▶ May lead to <b>deadlock</b></li></ul>

# Deadlock

- ▶ Deadlock is a situation where a set of processes are blocked.
- ▶ Each process in the set waits for some event from some other process in the same set.
- ▶ Example #1
  - ▶ A system has 2 disk drives
  - ▶  $P_1$  and  $P_2$  each hold one disk drive and each needs the other one
- ▶ Example #2
  - ▶ Semaphores  $A$  and  $B$ , initialized to 1

$P_0$	$P_1$
wait (A);	wait(B)
wait (B);	wait(A)



# Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- ▶ **Mutual exclusion:** only one process at a time can use a resource
- ▶ **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- ▶ **No preemption:** a resource can be released only voluntarily by the process holding it after that process has completed its task
- ▶ **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$

# Methods for Handling Deadlocks

## ▶ Prevention

- ▶ Ensure that the system will never enter a deadlock state

## ▶ Avoidance

- ▶ Ensure that the system will never enter an unsafe state

## ▶ Detection

- ▶ Allow the system to enter a deadlock state and then recover

# Deadlock Prevention

Restrain at least one of the causes of deadlock

- ▶ Mutual Exclusion – The mutual-exclusion condition must hold for non-sharable resources.
- ▶ Hold and Wait – we must guarantee that whenever a process requests a resource, it does not hold any other resources
- ▶ There is no guarantee that a file that is released by a process will be in the same state when it will get for the next time.

# Deadlock Prevention

## ► No Preemption

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are preempted.
- Preempted resources are added to the list of available resources.
- A process will be restarted only when it can regain its old resources, and the new ones that it is requesting



# Deadlock Prevention

## ► Circular Wait

- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration. For example:

1. R1
2. R2
3. R3
4. R4

If  $P_i$  holds R3 and requests for R2 then it has release R3 first and re-request for R2 then R3

**What if R3 is a file!!**

# Deadlock Avoidance and Safe State

- ▶ When a process requests an available resource, the system must decide if immediate allocation leaves the system in a ***safe state***
- ▶ A system is in a safe state only if there exists a ***safe sequence of execution***
- ▶ Total number of resources=12


Process	Max Needs	Allocated	Current Needs
P0	10	5	5
P1	4	2	2
P2	7	3	4

- ▶ P0 requests one more resource dynamically. Can it be granted?
  - ▶  $\langle P1, P2, P0 \rangle$  is a **safe sequence of execution**

# Safe State (continued)

- ▶ If a system is in safe state  $\Rightarrow$  no deadlocks
- ▶ If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- ▶ Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state

# Safe, Unsafe, Deadlock State



# Avoidance algorithms

- ▶ For a single instance of a resource type, use a *resource-allocation graph (RAG)*
- ▶ For multiple instances of a resource type, use the *banker's algorithm*


# RAG for Deadlock avoidance

- ▶  $G = (V, E)$
- ▶ V: Processes and Resources
- ▶ E: Request edge, Allocation edge, Claim edge

Process:




*Claim edge:*




Resource:



*Request edge:*




*Assignment edge:*



# RAG for Deadlock avoidance

- ▶ *Claim edge*  $P_i \text{-----} \rightarrow R_j$  indicates that process  $P_j$  may request resource  $R_j$ ; which is represented by a dashed line
- ▶ A claim edge converts to a request edge when a process **requests** a resource
- ▶ A request edge converts to an assignment edge when the resource is **allocated** to the process


# RAG for Deadlock avoidance








# RAG for Deadlock avoidance



# RAG for Deadlock avoidance

**Presence of Cycle in RAG may lead to DEADLOCK**

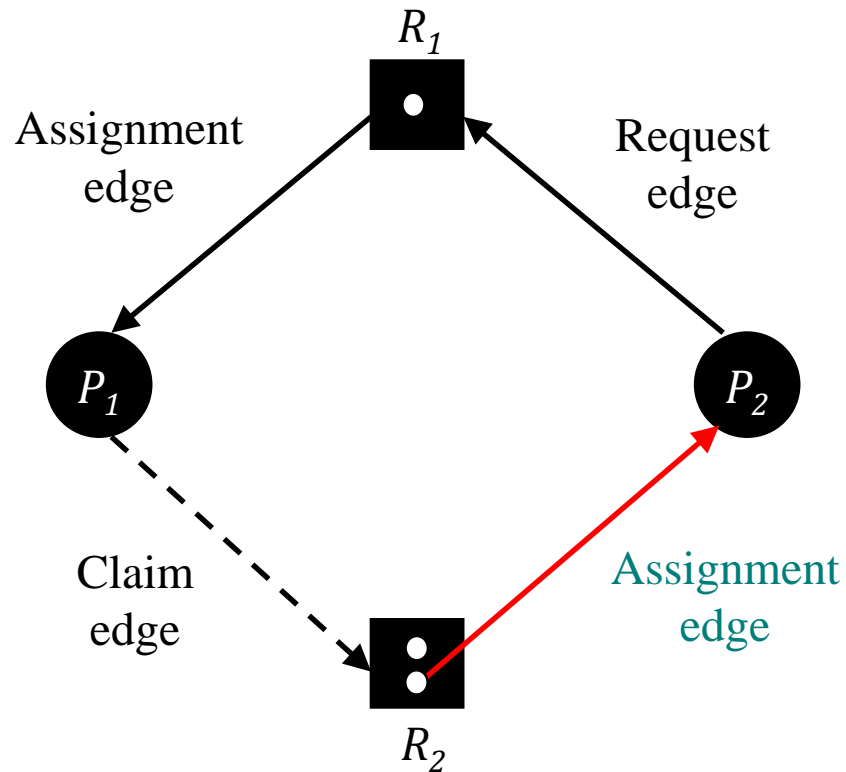


# RAG for Deadlock avoidance

- ▶ Suppose that process  $P_i$  requests a resource  $R_j$
- ▶ The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

# RAG for Deadlock avoidance

- RAG algorithm doesn't work for multiple instances of resources.



# Banker's Algorithm

- ▶ When a process makes a request for a set of resources,
- ▶ Resource Request Algorithm:
  - ▶ assume that the request is granted,
  - ▶ Update the system state accordingly,
  - ▶ Then, determine if the result is a safe state (Safety algo)
    - If so, grant the request
    - if not, block the process until it is safe to grant the request.

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- ▶ **Available**: Vector of length  $m$ . If **Available**  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- ▶ **Max**:  $n \times m$  matrix. If **Max**  $[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- ▶ **Allocation**:  $n \times m$  matrix. If **Allocation** $[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
- ▶ **Need**:  $n \times m$  matrix. If **Need** $[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$Need [i, j] = Max[i, j] - Allocation [i, j]$$

# Safety Algorithm

1. Initialize  $Work = Available$   
Initialize  $Finish[i] = False$ , for  $i = 1, 2, 3, \dots, n$
2. Find an  $i$  such that:  
 $Finish[i] == False$  and  $Need[i] \leq Work$   
If no such  $i$  exists, go to step 4.
3.  $Work = Work + Allocation[i]$   
 $Finish[i] = True$   
goto step 2
4. if  $Finish[i] == True$  for all  $i$ , then the system is in a safe state.

# Banker's Algorithm: Example

Initially Available=(10,5,7)

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	3	3	2
P1	2	0	0	1	2	2			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

P1 makes  $\text{Request}_1 = (1, 0, 2)$



# Banker's Algorithm: Example

Updated State.

Let's check for a safe sequence.

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	2	3	0
P1	3	0	2	0	2	0			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

P1 makes  $\text{Request}_1 = (1, 0, 2)$

# Banker's Algorithm: Example

P1 can be completed.

Pro cess	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	5	3	2
P1	0	0	0	0	0	0			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

P1 makes  $\text{Request}_1 = (1, 0, 2)$

# Banker's Algorithm: Example

P3 can be completed.

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	7	4	3
P1	0	0	0	0	0	0			
P2	3	0	2	6	0	0			
P3	0	0	0	0	0	0			
P4	0	0	2	4	3	1			

P1 makes  $\text{Request}_1 = (1, 0, 2)$

# Banker's Algorithm: Example

P4 can be completed.

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	7	4	5
P1	0	0	0	0	0	0			
P2	3	0	2	6	0	0			
P3	0	0	0	0	0	0			
P4	0	0	0	0	0	0			

P1 makes  $\text{Request}_1 = (1, 0, 2)$

# Banker's Algorithm: Example

P0 can be completed.

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	0	0	0	0	0	7	5	5
P1	0	0	0	0	0	0			
P2	3	0	2	6	0	0			
P3	0	0	0	0	0	0			
P4	0	0	0	0	0	0			

P1 makes  $\text{Request}_1 = (1, 0, 2)$

# Banker's Algorithm: Example

P2 can be completed.

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	0	0	0	0	0	10	5	7
P1	0	0	0	0	0	0			
P2	0	0	0	0	0	0			
P3	0	0	0	0	0	0			
P4	0	0	0	0	0	0			

P1 makes  $\text{Request}_1 = (1, 0, 2)$

# Banker's Algorithm: Example

Safe Sequence: <P1,P3,P4,P0,P2>

Pro cess	Allocation				Need				Available		
	A	B	C		A	B	C		A	B	C
P0	0	1	0		7	4	3		2	3	0
P1	3	0	2		0	2	0				
P2	3	0	2		6	0	0				
P3	2	1	1		0	1	1				
P4	0	0	2		4	3	1				

P1 makes  $\text{Request}_1 = (1, 0, 2) \rightarrow \text{GRANTED}$

# Banker's Algorithm: Example

Process	Allocation				Need				Available		
	A	B	C		A	B	C		A	B	C
P0	0	1	0		7	4	3		2	3	0
P1	3	0	2		0	2	0				
P2	3	0	2		6	0	0				
P3	2	1	1		0	1	1				
P4	0	0	2		4	3	1				

Now, P4 makes  $\text{Request}_4 = (3, 3, 0) \rightarrow \text{NOT AVAILABLE}$



# Banker's Algorithm: Example

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	2	3	0
P1	3	0	2	0	2	0			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

Now, P0 makes **Request<sub>0</sub> = (0, 2, 0)**

# Banker's Algorithm: Example

Updated State.

Let's check for a safe sequence.

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	3	0	7	2	3	2	1	0
P1	3	0	2	0	2	0			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

Now, P0 makes  $\text{Request}_0 = (0, 2, 0)$

# Banker's Algorithm: Example

No process can complete

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	3	0	7	2	3	2	1	0
P1	3	0	2	0	2	0			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

Now, P0 makes  $\text{Request}_0 = (0, 2, 0) \rightarrow$  NOT GRANTED

# Banker's Algorithm: Example

Revert back to previous state

Process	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	2	3	0
P1	3	0	2	0	2	0			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

Disadvantage: Time complexity is very high.

# Deadlock Detection and Recovery

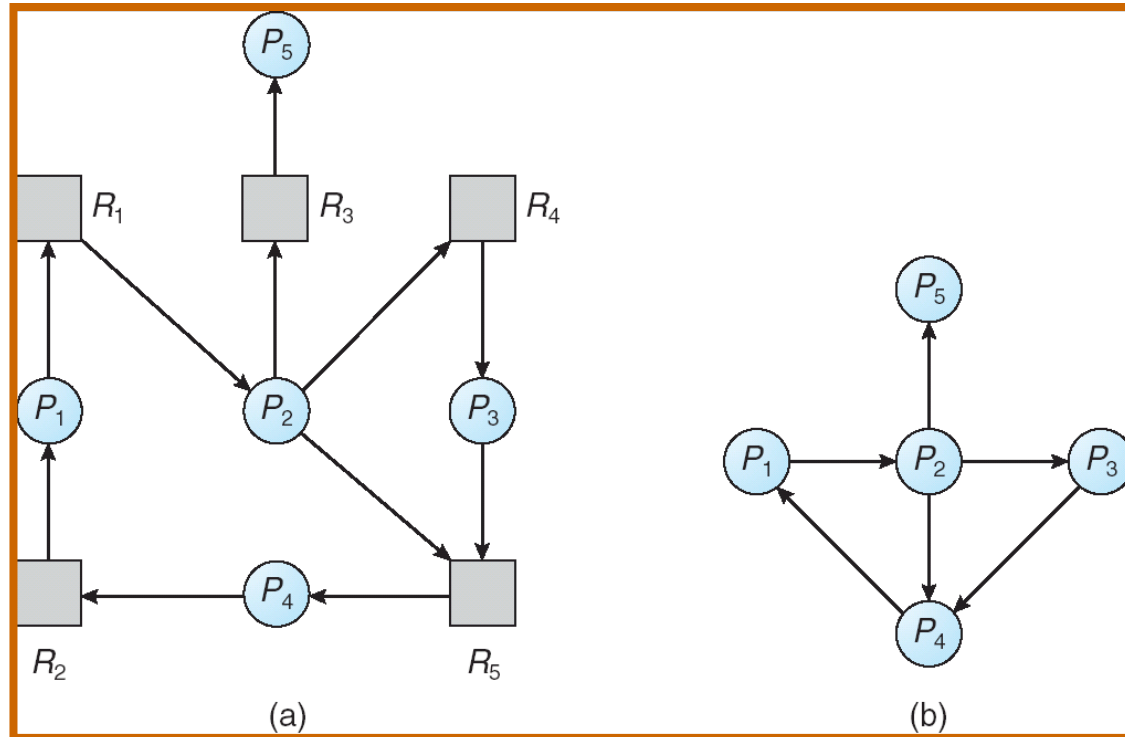
- ▶ An algorithm that examines the state of the system to detect whether a deadlock has occurred
- ▶ And an algorithm to recover from the deadlock

# Deadlock Detection:

## Single Instance of Each Resource Type

- ▶ Requires the creation and maintenance of a wait for graph (WFG)
  - ▶ variant of the resource-allocation graph
  - ▶ The graph is obtained by **removing** the resource nodes from a RAG and **collapsing** the appropriate edges
  - ▶ Consequently; all nodes are processes
  - ▶  $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- ▶ Periodically invoke an algorithm that searches for a cycle in the graph
  - ▶ If there is a cycle, there exists a deadlock
  - ▶ An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph

# Deadlock Detection: Single Instance of Each Resource Type



Resource-Allocation Graph

Corresponding wait-for graph

# Deadlock Detection: Multiple Instances of Resource Types

## Variation of Banker's Algorithm

1. Initialize  $Work = Available$   
Initialize  $Finish[i] = False$ , for  $i = 1, 2, 3, \dots, n$
2. Find an  $i$  such that:  
 $Finish[i] == False$  and  $Request[i] \leq Work$   
If no such  $i$  exists, go to step 4.
3.  $Work = Work + Allocation[i]$   
 $Finish[i] = True$   
goto step 2
4. if  $Finish[i] == true$  for all  $i$ , then the system is not in  
deadlocked state



# Deadlock Detection Algorithm: Example

Sequence  $\langle P0, P2, P3, P1, P4 \rangle$  results in  $Finish[i] == true$  for all  $i$   
 $\Rightarrow$  *System is not in deadlocked state*

Pro cess	Allocation				Request				Available		
	A	B	C		A	B	C		A	B	C
P0	0	1	0		0	0	0		0	0	0
P1	2	0	0		2	0	2				
P2	3	0	3		0	0	0				
P3	2	1	1		1	0	0				
P4	0	0	2		0	0	2				

# Deadlock Detection Algorithm: Example

Finish[0] is True but **P1, P2, P3, P4** are deadlocked.

Process	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	0	0	0	0	0	0
P1	2	0	0	2	0	2			
P2	3	0	3	0	0	1			
P3	2	1	1	1	0	0			
P4	0	0	2	0	0	2			

# Recovery from Deadlock

- ▶ Two Approaches
  - ▶ Process termination
  - ▶ Resource preemption

# Recovery from Deadlock: Process Termination

- ▶ Abort all deadlocked processes
  - ▶ This approach will break the deadlock, but at great expense
- ▶ Abort one process at a time until the deadlock cycle is eliminated
  - ▶ This approach incurs considerable overhead, since, after each process is aborted, a deadlock-detection algorithm must be re-invoked to determine whether any processes are still deadlocked
- ▶ Many factors may affect which process is chosen for termination
  - ▶ What is the priority of the process?
  - ▶ How long has the process run so far and how much longer will the process need to run before completing its task?
  - ▶ How many and what type of resources has the process used?
  - ▶ How many more resources does the process need in order to finish its task?
  - ▶ How many processes will need to be terminated?

# Recovery from Deadlock: Resource Preemption

- ▶ With this approach, we successively preempt some resources from processes and give these resources to other processes until the deadlock cycle is broken
- ▶ When preemption is required to deal with deadlocks, then three issues need to be addressed:
  - ▶ **Selecting a victim** – Which resources and which processes are to be preempted?
  - ▶ **Rollback** – If we preempt a resource from a process, what should be done with that process?
  - ▶ **Starvation** – How do we ensure that starvation will not occur? That is, how can we guarantee that resources will not always be preempted from the same process?

# Summary

- ▶ Both deadlock avoidance and deadlock detection and recovery are expensive
- ▶ **Do nothing:** ignore the problem altogether and pretend that deadlocks never occur in the system (used by Windows and Unix)
  - ▶ System administrator/user will restart the processes